

## Linear Regression Assignment

In a watershed monitoring project, automated samplers are often used to collect water samples. However, these automated samplers may fail from time to time due to various instrumentation issues, such as power failure and damaged instruments. When the automated sampler failed to collect water samples, it leaves a data gap from a certain period of time.

Base flow nutrient concentration data are often linearly interpolated (or sometimes extrapolated) using nutrient data of the last and next sample collected. This approach cannot be applied to event flow nutrient and sediment (hereinafter refer as nutrient) data, especially phosphorus (P) and sediment (TSS), because of the large variations in nutrient concentration that can be caused by flow.

The objective of this task are to:

- 1) develop a suitable model using existing dissolved reactive P (DRP), total P (TP), and TSS
- 2) predict the missing DRP, TP, TSS data using the model developed in (1)

```
In [99]: import pandas as pd
import numpy as np
import seaborn as sns
import statsmodels.formula.api as smf
import matplotlib.pyplot as plt
from scipy import stats
```

```
In [21]: # import data from excel
data = pd.read_excel('JY_linear_reg_data.xlsx', sheet_name = 'Data')
data.head()
```

Out[21]:

	Sample date	Site	Flow (cms)	DRP (mg P/L)	TP (mg P/L)	TSS (mg/L)	VSS (mg/L)
0	2015-05-19	S11	0.028210	0.018	0.0015	22.333667	19.264296
1	2015-06-16	S11	0.189993	0.027	0.1770	1006.666667	892.000000
2	2015-06-30	S11	1.296460	0.203	1.2590	1338.666667	1194.666667
3	2015-08-18	S11	0.131722	0.050	1.3280	2026.000000	1812.000000
4	2015-08-25	S11	0.120204	0.010	3.1180	1722.000000	1560.000000

```
In [22]: # drop 'VSS' column, rename the columns into shorter names

data = data.drop(columns=['VSS (mg/L)'])
data = data.rename(columns={'Sample date': 'Date', 'Flow (cms)': 'Flow', 'DRP (mg P/L)': 'DRP', 'TP (mg P/L)': 'TP', 'TSS (mg/L)': 'TSS'})
data.head()
```

Out[22]:

	Date	Site	Flow	DRP	TP	TSS
0	2015-05-19	S11	0.028210	0.018	0.0015	22.333667
1	2015-06-16	S11	0.189993	0.027	0.1770	1006.666667
2	2015-06-30	S11	1.296460	0.203	1.2590	1338.666667
3	2015-08-18	S11	0.131722	0.050	1.3280	2026.000000
4	2015-08-25	S11	0.120204	0.010	3.1180	1722.000000

```
In [117]: data_11 = data[data['Site'] == 'S11']
data_12 = data[data['Site'] == 'S12']
print(data_11.tail())
print(data_12.tail())
```

	Date	Site	Flow	DRP	TP	TSS
24	2018-09-11	S11	0.120500	0.052	0.338	363.333333
25	2018-09-11	S11	0.115396	0.014	0.498	162.000000
26	2018-09-26	S11	0.085015	0.010	0.293	474.500000
27	2018-10-11	S11	0.113744	0.047	0.190	148.000000
28	2018-10-11	S11	0.156634	NaN	NaN	NaN
	Date	Site	Flow	DRP	TP	TSS
60	2018-08-31	S12	0.028923	NaN	NaN	NaN
61	2018-09-06	S12	0.024730	NaN	NaN	NaN
62	2018-09-26	S12	0.022594	NaN	NaN	NaN
63	2018-10-11	S12	0.024522	0.001	0.042	11.666667
64	2018-10-11	S12	0.044039	0.003	0.040	20.333333

```
In [43]: # check the number of null values
print('S11')
print(data_11.isnull().sum())
print('S12')
print(data_12.isnull().sum())
```

```
S11
Date      0
Site      0
Flow      0
DRP       3
TP        3
TSS       3
dtype: int64
S12
Date      0
Site      0
Flow      0
DRP       4
TP        4
TSS       4
dtype: int64
```

```
In [47]: # create separate dataframe that only contain null values
data_11_null = data_11[data_11.isnull().any(axis=1)]
data_12_null = data_12[data_12.isnull().any(axis=1)]
print('S11')
print(data_11_null)
print('S12')
print(data_12_null)
```


```
S11
      Date Site      Flow  DRP  TP  TSS
21 2018-06-22 S11  0.103414  NaN  NaN  NaN
22 2018-06-27 S11  0.189099  NaN  NaN  NaN
28 2018-10-11 S11  0.156634  NaN  NaN  NaN
S12
      Date Site      Flow  DRP  TP  TSS
59 2018-08-21 S12  0.128687  NaN  NaN  NaN
60 2018-08-31 S12  0.028923  NaN  NaN  NaN
61 2018-09-06 S12  0.024730  NaN  NaN  NaN
62 2018-09-26 S12  0.022594  NaN  NaN  NaN
```

```
In [195]: # drop the null values
data_11 = data_11.dropna()
data_12 = data_12.dropna()
print('S11')
print(data_11.isnull().sum())
print('S12')
print(data_12.isnull().sum())
```

```
S11
Date      0
Site      0
Flow      0
DRP       0
TP        0
TSS       0
dtype: int64
S12
Date      0
Site      0
Flow      0
DRP       0
TP        0
TSS       0
dtype: int64
```

```
In [50]: # visualize the relationship between flow and responses
sns.pairplot(data_11, x_vars=['DRP', 'TP', 'TSS'], y_vars=['Flow'], height=6, aspect=0.7)
sns.pairplot(data_12, x_vars=['DRP', 'TP', 'TSS'], y_vars=['Flow'], height=6, aspect=0.7)
```

```
Out[50]: <seaborn.axisgrid.PairGrid at 0x175d4795240>
```



## Objective 1

Develop a suitable model using existing dissolved reactive P (DRP), total P (TP), and TSS

```
In [80]: list_of_comparisons = ['DRP ~ Flow', 'TP ~ Flow', 'TSS ~ Flow']

for x in list_of_comparisons:
    print(x)
    result = smf.ols(formula=x, data=data_11).fit()
    print(result.summary())
```

DRP ~ Flow

## OLS Regression Results

```

=====
Dep. Variable:          DRP      R-squared:                0.719
Model:                  OLS      Adj. R-squared:           0.708
Method:                 Least Squares      F-statistic:        61.48
Date:                  Fri, 04 Oct 2019      Prob (F-statistic):    4.50e-08
Time:                  11:01:56      Log-Likelihood:       61.156
No. Observations:      26      AIC:                 -118.3
Df Residuals:          24      BIC:                 -115.8
Df Model:              1
Covariance Type:       nonrobust
=====

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept    0.0068      0.006       1.199     0.242     -0.005      0.018
Flow         0.1487      0.019       7.841     0.000      0.110      0.188
=====

```

```

=====
Omnibus:            13.007      Durbin-Watson:           2.306
Prob(Omnibus):      0.001      Jarque-Bera (JB):        11.946
Skew:               1.467      Prob(JB):                0.00255
Kurtosis:           4.554      Cond. No.                4.15
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

TP ~ Flow

## OLS Regression Results

```

=====
Dep. Variable:          TP      R-squared:                0.036
Model:                  OLS      Adj. R-squared:          -0.005
Method:                 Least Squares      F-statistic:           0.8870
Date:                  Fri, 04 Oct 2019      Prob (F-statistic):    0.356
Time:                  11:01:56      Log-Likelihood:       -26.991
No. Observations:      26      AIC:                 57.98
Df Residuals:          24      BIC:                 60.50
Df Model:              1
Covariance Type:       nonrobust
=====

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept    0.4394      0.167       2.626     0.015      0.094      0.785
Flow         0.5301      0.563       0.942     0.356     -0.632      1.692
=====

```

```

=====
Omnibus:            29.806      Durbin-Watson:           1.758
Prob(Omnibus):      0.000      Jarque-Bera (JB):        59.814
Skew:               2.346      Prob(JB):                1.03e-13
Kurtosis:           8.762      Cond. No.                4.15
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

TSS ~ Flow

## OLS Regression Results

```

=====
Dep. Variable:          TSS      R-squared:                0.003
Model:                  OLS      Adj. R-squared:          -0.038
Method:                 Least Squares      F-statistic:           0.08284
Date:                  Fri, 04 Oct 2019      Prob (F-statistic):    0.776
Time:                  11:01:56      Log-Likelihood:       -236.65
No. Observations:      26      AIC:                 477.3
Df Residuals:          24      BIC:                 479.8
=====

```

Df Model: 1  
 Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
Intercept	813.7883	531.769	1.530	0.139	-283.730	1911.306
Flow	514.7888	1788.619	0.288	0.776	-3176.739	4206.316
Omnibus:	55.816	Durbin-Watson:	1.639			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	362.575			
Skew:	4.085	Prob(JB):	1.85e-79			
Kurtosis:	19.369	Cond. No.	4.15			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [52]: list_of_comparisons = ['DRP ~ Flow', 'TP ~ Flow', 'TSS ~ Flow']

for x in list_of_comparisons:
    print(x)
    result = smf.ols(formula=x, data=data_12).fit()
    print(result.summary())
```



DRP ~ Flow

## OLS Regression Results

```

=====
Dep. Variable:          DRP      R-squared:                0.413
Model:                  OLS      Adj. R-squared:           0.393
Method:                 Least Squares      F-statistic:        21.10
Date:                   Fri, 04 Oct 2019    Prob (F-statistic):    7.32e-05
Time:                   10:13:48          Log-Likelihood:       47.461
No. Observations:       32              AIC:                 -90.92
Df Residuals:           30              BIC:                 -87.99
Df Model:                1
Covariance Type:        nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept      0.0116      0.013      0.918      0.366      -0.014      0.037
Flow           0.4836      0.105      4.593      0.000      0.269      0.699
=====
Omnibus:                47.570    Durbin-Watson:           1.320
Prob(Omnibus):           0.000    Jarque-Bera (JB):        236.046
Skew:                    3.113    Prob(JB):                5.54e-52
Kurtosis:                14.759    Cond. No.                 10.6
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

TP ~ Flow

## OLS Regression Results

```

=====
Dep. Variable:          TP      R-squared:                0.558
Model:                  OLS      Adj. R-squared:           0.543
Method:                 Least Squares      F-statistic:        37.81
Date:                   Fri, 04 Oct 2019    Prob (F-statistic):    9.20e-07
Time:                   10:13:48          Log-Likelihood:       0.093905
No. Observations:       32              AIC:                 3.812
Df Residuals:           30              BIC:                 6.744
Df Model:                1
Covariance Type:        nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept      0.0357      0.055      0.644      0.524      -0.078      0.149
Flow           2.8447      0.463      6.149      0.000      1.900      3.790
=====
Omnibus:                13.785    Durbin-Watson:           2.480
Prob(Omnibus):           0.001    Jarque-Bera (JB):        15.234
Skew:                    1.211    Prob(JB):                0.000492
Kurtosis:                5.359    Cond. No.                 10.6
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

TSS ~ Flow

## OLS Regression Results

```

=====
Dep. Variable:          TSS      R-squared:                0.556
Model:                  OLS      Adj. R-squared:           0.541
Method:                 Least Squares      F-statistic:        37.55
Date:                   Fri, 04 Oct 2019    Prob (F-statistic):    9.76e-07
Time:                   10:13:48          Log-Likelihood:      -230.79
No. Observations:       32              AIC:                 465.6
Df Residuals:           30              BIC:                 468.5
=====

```

```

Df Model: 1
Covariance Type: nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept    -143.3081      75.413     -1.900      0.067     -297.323      10.706
Flow          3854.9033     629.085      6.128      0.000     2570.141     5139.665
=====
Omnibus:                20.257   Durbin-Watson:                2.516
Prob(Omnibus):           0.000   Jarque-Bera (JB):           68.576
Skew:                    0.986   Prob(JB):                   1.29e-15
Kurtosis:                9.895   Cond. No.                    10.6
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## Discussion for objective 1

The quality of the OLS models was evaluated based on p-value and R2:

- 1) S11-DRP - will work well
- 2) S11-TP - will not work well
- 3) S11-TSS - will not work well based on R2, p-value, and high Y-intercept value
- 4) S12-DRP - will work well
- 5) S12-TP - will work well
- 6) S12-TSS - appear that it will work well based on R2 and p-value, but the Y-intercept is largely negative

## Objective 2

Predict the missing DRP, TP, TSS data using the model developed in (1)

```

In [176]: # predict values and export the values into a dataframe for S11
list_of_comparisons = ['DRP ~ Flow', 'TP ~ Flow', 'TSS ~ Flow']
predictions_11 = pd.DataFrame(columns=list(list_of_comparisons))

for x in list_of_comparisons:
    result = smf.ols(formula=x, data=data_11).fit()
    prednull_11 = result.predict(data_11_null['Flow'])
    predictions_11.loc[:, x] = prednull_11
predictions_11

```

Out[176]:

	DRP ~ Flow	TP ~ Flow	TSS ~ Flow
21	0.022142	0.494259	867.024807
22	0.034885	0.539682	911.134534
28	0.030057	0.522472	894.421904




```
In [199]: # add the values back into this new dataframe - getting the dataframe ready to merge with
           the main dataframe
           predictions_11 ['Flow'] = data_11_null ['Flow']
           predictions_11 ['Date'] = data_11_null ['Date']
           predictions_11 ['Site'] = 'S11'
           predictions_11 = predictions_11.rename(columns={'DRP ~ Flow': 'DRP', 'TP ~ Flow': 'TP', 'T
           SS ~ Flow': 'TSS'})
           predictions_11
```

Out[199]:

	DRP	TP	TSS	Flow	Date	Site
21	0.022142	0.494259	867.024807	0.103414	2018-06-22	S11
22	0.034885	0.539682	911.134534	0.189099	2018-06-27	S11
28	0.030057	0.522472	894.421904	0.156634	2018-10-11	S11

```
In [219]: analytes = ['DRP', 'TP', 'TSS']

for x in analytes:
    plt.scatter(data_11['Flow'], data_11[x], label='actual')
    plt.scatter(predictions_11['Flow'], predictions_11[x], label='predicted')
    plt.legend()
    plt.xlabel('Flow')
    plt.ylabel(x)
    plt.show()
```



```
In [220]: data_11.shape
```

```
Out[220]: (26, 6)
```

```
In [221]: merged_11 = pd.concat([data_11, predictions_11], sort=True)
merged_11 = merged_11.sort_values(by=['Date'])
merged_11.shape
```

Out[221]: (29, 6)

```
In [223]: # predict values and export the values into a dataframe for S12
list_of_comparisons = ['DRP ~ Flow', 'TP ~ Flow', 'TSS ~ Flow']
predictions_12 = pd.DataFrame(columns=list(list_of_comparisons))




for x in list_of_comparisons:
    result = smf.ols(formula=x, data=data_12).fit()
    prednull_12 = result.predict(data_12_null['Flow'])
    predictions_12.loc[:, x] = prednull_12

# add the values back into this new dataframe - getting the dataframe ready to merge with
# the main dataframe
predictions_12 ['Flow'] = data_12_null ['Flow']
predictions_12 ['Date'] = data_12_null ['Date']
predictions_12 ['Site'] = 'S12'
predictions_12 = predictions_12.rename(columns={'DRP ~ Flow': 'DRP', 'TP ~ Flow': 'TP', 'TSS ~ Flow': 'TSS'})
predictions_12.shape
```

Out[223]: (4, 6)

```
In [222]: analytes = ['DRP', 'TP', 'TSS']

for x in analytes:
    plt.scatter(data_12['Flow'], data_12[x], label='actual')
    plt.scatter(predictions_12['Flow'], predictions_12[x], label='predicted')
    plt.legend()
    plt.xlabel('Flow')
    plt.ylabel(x)
    plt.show()
```



```
In [225]: merged_12 = pd.concat([data_12, predictions_12], sort=True)
merged_12 = merged_12.sort_values(by=['Date'])
merged_12.shape
```

Out[225]: (36, 6)

```
In [229]: # merge all the data back together
final_merged = pd.concat([merged_11, merged_12])
print(final_merged.shape)
print(final_merged.isnull().sum())
```

```
(65, 6)
DRP      0
Date      0
Flow      0
Site      0
TP        0
TSS       0
dtype: int64
```

## Discussion for objective 2

From the figures above where predicted and actual values were plotted:

- 1) S11-DRP - worked well
- 2) S11-TP - worked okay (within somewhat expected range, but the model had weak R2 and p-value)
- 3) S11-TSS - worked okay (within somewhat expected range)
- 4) S12- DRP - worked well
- 5) S12 - TP - worked well
- 6) S12 - TSS - did not work at all (some of the predicted TSS concentrations were negative)

In summary, this OLS model worked well for DRP and TP, but not TSS. Another suitable model is needed to predict TSS concentration. A potential approach is to multiple-linear-regression to incorporate several parameters such as precipitation data (antecedent condition) and time of the year (seasonality - land cover density).

In [ ]: